



AI SATURDAYS LAGOS

Week 4: Relational Data

Instructor: 'Tayo Jabar



Outline

1. Relational Data
2. Entity Relationships
3. Manipulating Relational Data
(Pandas & SQLite)
4. Joins






Relational Data

1. Overview
2. Primary & Foreign Keys
3. Indexes

Overview of Relational Data

Key word: Relation

- Organized, inter-related data
- Tabular structure
 - Rows
 - Columns
- Row: Record, Tuple
- Column: Feature, Attribute



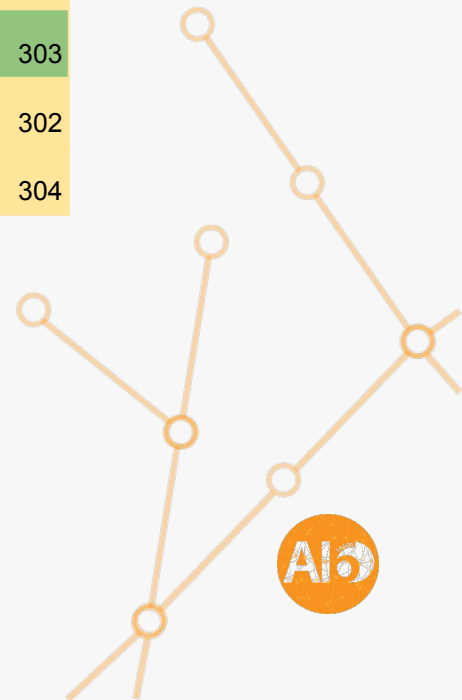
Relational Data (...cont'd)

ID	First_Name	Last_Name	Gender	Email	Profession
101	Tayo	Jabar	Male	a@b.com	301
102	Kenechi	Dukor	Male	b@c.com	
103	Tejumade	Afonja	Female	c@d.com	303
104	Fola	Animashaun	Female	d@e.com	302
105	Zazzu	Olalomi	Male	e@f.com	304

Legend

	Row
	Column
	Primary Key

A **Primary Key** is a column/feature that is unique for every record in a table.



Primary & Foreign Keys

Customers

ID	First_Name	Last_Name	Gender	Email	Profession
1	Tayo	Jabar	Male	a@b.com	301
2	Kenechi	Dukor	Male	b@c.com	302
3	Tejumade	Afonja	Female	c@d.com	303
4	Fola	Animashaun	Female	d@e.com	302
5	Zazzu	Olalomi	Male	e@f.com	304

Professions

ID	Profession_Name
301	Doctor
302	Engineer
303	Professor
304	Mechanic

A **Foreign Key** is a primary key in one table that is used as an identifier in another table.

The ID column in the Professions table is a **Primary Key** in that table.

The Professions column in the Customers table is a **Foreign Key**, identifying the professions in the Customers table.

Indexes

Location	ID	First_Name	Last_Name	Gender	Email	Profession
100	1	Tayo	Jabar	Male	a@b.com	301
200	2	Kenechi	Dukor	Male	b@c.com	302
300	3	Tejumade	Afonja	Female	c@d.com	303
400	4	Fola	Animashaun	Female	d@e.com	302
500	5	Zazzu	Olalomi	Male	e@f.com	304

Last Name	Location
Afonja	300
Animashaun	400
Dukor	200
Jabar	100
Olalomi	500

An **Index** is a database structure used to help speed up searching for an item in a database table.

An **Index Table** is a separate “sorted table” that holds the location of records within a table.

Quiz

Question A: What is a relational database?

- a) A database that plays videos
- b) A database that uses tables to store and organize data
- c) A database that only works on mobile devices
- d) A database for music streaming

Question B: What does the primary key in a relational database table represent?

- a) The main table in the database
- b) A unique identifier for each row in a table
- c) A way to delete data from a table
- d) A foreign key in another table





Entity Relationships

1. One to One
2. One to Many
3. Many to Many



One to One

- This shows an exclusive/strict case of one record in one table only matching one record in a second table.



Customers

ID	First_Name	Last_Name	Gender	Email	Profession
1	Tayo	Jabar	Male	a@b.com	301
2	Kenechi	Dukor	Male	b@c.com	302
3	Tejumade	Afonja	Female	c@d.com	303
4	Fola	Animashaun	Female	d@e.com	302
5	Zazzu	Olalomi	Male	e@f.com	304

Username

Customer ID	Username
1	tayo_jabar
2	kenny_dukor
3	teju_afonja
4	folo_animashaun
5	zazzu_olalomi

One to Many / Many to One

- This illustrates a case where one record in one table could match multiple records in another table or vice-versa.



Professions

ID	Profession_Name
301	Doctor
302	Engineer
303	Professor
304	Mechanic

Customers

ID	First_Name	Last_Name	Gender	Email	Profession
1	Tayo	Jabar	Male	a@b.com	301
2	Kenechi	Dukor	Male	b@c.com	302
3	Tejumade	Afonja	Female	c@d.com	303
4	Fola	Animashaun	Female	d@e.com	302
5	Zazzu	Olalomi	Male	e@f.com	304

Many to Many

- This illustrates a case where multiple records in one table could match multiple records in another table or vice-versa.



Customers


ID	First_Name	Last_Name	Gender	Email	Profession
1	Tayo	Jabar	Male	a@b.com	301
2	Kenechi	Dukor	Male	b@c.com	302
3	Tejumade	Afonja	Female	c@d.com	303
4	Fola	Animashaun	Female	d@e.com	302
5	Zazzu	Olalomi	Male	e@f.com	304

Items

ID	Items
531	Tissues
532	Samosa
533	Pens
534	Airpods



Many to Many: Associative Tables



ID	First_Name	Last_Name	Gender	Email	Profession
1	Tayo	Jabar	Male	a@b.com	301
2	Kenechi	Dukor	Male	b@c.com	302
3	Tejumade	Afonja	Female	c@d.com	303
4	Fola	Animashaun	Female	d@e.com	302
5	Zazzu	Olalomi	Male	e@f.com	304

Associative Table:
Items Bought / Sales

ID	Customer	Items
1101	101	531
1102	102	531
1103	103	533
1104	103	534
1105	104	532
1106	105	533
1107	102	534

ID	Items
531	Tissues
532	Samosa
533	Pens
534	Airpods



Quiz

Question A: What does an entity represent in entity-relationship modeling?

- a) A computer program
- b) A real-world object or concept
- c) A mathematical equation
- d) A data type in Python

Question B: Which of the following best describes a foreign key in entity-relationship modeling?

- a) A key that unlocks a treasure chest
- b) A key used to access the internet
- c) A key used to establish relationships between entities
- d) A key that opens a physical door




Quiz

Question C: In a many-to-many relationship, what is typically used to represent the relationship in a relational database?

- a) A foreign key in one of the related tables
- b) A new table that connects the two entities
- c) A primary key in each of the related tables
- d) A special "many-to-many" keyword

Question D: What is an associative entity in an entity-relationship diagram?

- a) An entity that represents a physical object
- b) An entity with a weak relationship to other entities
- c) An entity that connects two or more other entities
- d) An entity with no attributes





Manipulating Relational Data

1. Pandas
2. SQLite

Pandas

```
import pandas as pd
```

```
df = pd.DataFrame([(1, 'Kolter', 'Zico'),  
                  (2, 'Xi', 'Edgar'),  
                  (3, 'Lee', 'Mark'),  
                  (4, 'Mani', 'Shouvik'),  
                  (5, 'Gates', 'Bill'),  
                  (6, 'Musk', 'Elon')],  
                  columns=["id", "last_name", "first_name"])
```

df

	id	last_name	first_name
0	1	Kolter	Zico
1	2	Xi	Edgar
2	3	Lee	Mark
3	4	Mani	Shouvik
4	5	Gates	Bill
5	6	Musk	Elon



```
print (df.loc[1, "last_name"])
print ("#####")
print (df.loc[:, "last_name"])
print ("#####")
print (df.loc[:, ["last_name"]])
print ("#####")
```

```
print (df.iloc[4,1])
print (df.loc[4, "last_name"])
```

✓ 0.0s

Kolter

#####

id

1 Kolter

2 Xi

3 Lee

4 Mani

5 Gates

6 Musk

Name: last_name, dtype: object

#####

last_name

id

1 Kolter

2 Xi

3 Lee

4 Mani

5 Gates

6 Musk

#####

Bill

Mani

df.loc: based upon the "index" (i.e., effectively primary key) of the data frame.

df.iloc: based on the 0-indexed positional counter. 0-indexes can be used for both rows and columns.



SQLite

- A relational database management system (RDBMS)
- Direct connection to a database file. It is li(gh)te and is present on the edge
- Structured Query Language (SQL)
- Transactions maintain the ACID properties
 - Atomicity
 - Consistency
 - Isolation
 - Durability



SQLite (...cont'd)


```
import sqlite3
conn = sqlite3.connect("database.db")
cursor = conn.cursor()

### when you are done, call conn.close()

cursor.execute("""
CREATE TABLE person (
    id INTEGER PRIMARY KEY,
    last_name TEXT,
    first_name TEXT
);""")

cursor.execute("""
CREATE TABLE grades (
    person_id INTEGER PRIMARY KEY,
    hw1_grade INTEGER,
    hw2_grade INTEGER
);""")

conn.commit()
```



```
cursor.execute("INSERT INTO person VALUES (1, 'Kolter', 'Zico');")
cursor.execute("INSERT INTO person VALUES (2, 'Xi', 'Edgar');")
cursor.execute("INSERT INTO person VALUES (3, 'Lee', 'Mark');")
cursor.execute("INSERT INTO person VALUES (4, 'Mani', 'Shouvik');")
cursor.execute("INSERT INTO person VALUES (5, 'Gates', 'Bill');")
cursor.execute("INSERT INTO person VALUES (6, 'Musk', 'Elon');")

cursor.execute("INSERT INTO grades VALUES (5, 85, 95);")
cursor.execute("INSERT INTO grades VALUES (6, 80, 60);")
cursor.execute("INSERT INTO grades VALUES (100, 100, 100);")
```

Reading from an SQLite table

`SELECT <columns> FROM <tables> WHERE <conditions>`

```
✓ for row in cursor.execute("SELECT * FROM person;"):
    print(row)
```

```
(1, 'Kolter', 'Zico')
(2, 'Xi', 'Edgar')
(3, 'Lee', 'Mark')
(4, 'Mani', 'Shouvik')
(5, 'Gates', 'Bill')
(6, 'Musk', 'Elon')
```

```
pd.read_sql_query("SELECT * from person;", conn, index_col="id")
```

	last_name	first_name
id		
1	Kolter	Zico
2	Xi	Edgar
3	Lee	Mark
4	Mani	Shouvik
5	Gates	Bill
6	Musk	Elon

We Value Your
Feedback.

https://docs.google.com/forms/d/e/1FAIpQLSfoYXRTHJfiQo1Xi2PrgoNLKWOAUeiVl4MLrRrOACtYFzM8cg/viewform?usp=pp_url






Joins

1. Inner Join
2. Left Join
3. Right Join
4. Outer Join




SQL Joins


JOINS




Inner join



Full outer join




Left outer join



Right outer join

Image from [Numpy Ninja](#)



Given the following tables:

Person

	last_name	first_name
id		
1	Kolter	Zico
2	Xi	Edgar
3	Lee	Mark
4	Mani	Shouvik
5	Gates	Bill
6	Musk	Elon

Grades

	hw1_grade	hw2_grade
person_id		
5	85	95
6	80	60
100	100	100



Inner Join

Using Pandas

```
df_person = pd.read_sql_query("SELECT * FROM person", conn)
df_grades = pd.read_sql_query("SELECT * FROM grades", conn)
df_person.merge(df_grades, how="inner", left_on="id", right_on="person_id")
```

	id	last_name	first_name	person_id	hw1_grade	hw2_grade
0	5	Gates	Bill	5	85	95
1	6	Musk	Elon	6	80	60

Using SQL

```
pd.read_sql_query("SELECT * FROM person, grades WHERE person.id = grades.person_id", conn)
```

	id	last_name	first_name	person_id	hw1_grade	hw2_grade
0	5	Gates	Bill	5	85	95
1	6	Musk	Elon	6	80	60

Left Join

Using Pandas

```
df_person = pd.read_sql_query("SELECT * FROM person", conn)
df_grades = pd.read_sql_query("SELECT * FROM grades", conn)
df_person.merge(df_grades, how="left", left_on = "id", right_on="person_id")
```

	id	last_name	first_name	person_id	hw1_grade	hw2_grade
0	1	Kolter	Zico	NaN	NaN	NaN
1	2	Xi	Edgar	NaN	NaN	NaN
2	3	Lee	Mark	NaN	NaN	NaN
3	4	Mani	Shouvik	NaN	NaN	NaN
4	5	Gates	Bill	5.0	85.0	95.0
5	6	Musk	Elon	6.0	80.0	60.0

Using SQL

```
pd.read_sql_query("SELECT * FROM person LEFT JOIN grades ON person.id = grades.person_id" , conn)
```

Right Join

Using Pandas

```
df_person.merge(df_grades, how="right", left_on = "id", right_on="person_id")
```

	id	last_name	first_name	person_id	hw1_grade	hw2_grade
0	5.0	Gates	Bill	5	85	95
1	6.0	Musk	Elon	6	80	60
2	NaN	NaN	NaN	100	100	100

Using SQL

```
pd.read_sql_query("SELECT * FROM grades LEFT JOIN person ON grades.person_id = person.id" , conn)
```

✓ 0.0s

	person_id	hw1_grade	hw2_grade	id	last_name	first_name
0	5	85	95	5.0	Gates	Bill
1	6	80	60	6.0	Musk	Elon
2	100	100	100	NaN	None	None



Outer Join

```
df_person.merge(df_grades, how="outer", left_on = "id", right_on="person_id")
```

	id	last_name	first_name	person_id	hw1_grade	hw2_grade
0	1.0	Kolter	Zico	NaN	NaN	NaN
1	2.0	Xi	Edgar	NaN	NaN	NaN
2	3.0	Lee	Mark	NaN	NaN	NaN
3	4.0	Mani	Shouvik	NaN	NaN	NaN
4	5.0	Gates	Bill	5.0	85.0	95.0
5	6.0	Musk	Elon	6.0	80.0	60.0
6	NaN	NaN	NaN	100.0	100.0	100.0

Quiz

Question A: In an SQL LEFT JOIN, which table's data is preserved entirely?

- a) Left table
- b) Right table
- c) Both tables
- d) None of the tables

Question B: What is the primary criterion for joining tables in a SQL query?

- a) The number of rows in each table
- b) The data types of the columns
- c) A common column or key between the tables
- d) The order of the columns



Thank You



Resources and External Links

- Practical Data Science: Relational Data - CMU AI, Zico Kolter, Pat Virtue
 - https://www.datasciencecourse.org/slides/15388_S22_Lecture_4_relational_data.pdf
 - http://www.datasciencecourse.org/notes/relational_data
- What is a Relational Database? - IBM, [Relational Databases](#)
- SQLite, <https://www.sqlite.org/index.html>

